



# ТЕХНОСФЕРА

## Лекция Online advertising 2

Владимир Гулин

10 апреля 2020 г.

# План лекции

Kaggle ads click challenges

SOTA

Practical Lessons from industry companies

## Kaggle ads click challenges



# Kaggle ads click challenges

## Criteo leaderboard

1	—	3 Idiots	  	0.44463	13	4y
2	—	Michael Jahrer and Jeong-Yoo...	 	0.44527	61	4y
3	—	beile		0.44610	67	4y

## Avazu leaderboard

1	—	4 Idiots	   	0.3791384	273	3y
2	—	Owen		0.3803652	94	3y
3	—	Random Walker	 	0.3806351	242	3y

SOTA

# Linear Prediction Models

$$\hat{y} = f(\mathbf{w}^T \mathbf{x})$$

- Pros
  - Highly efficient and scalable
  - Explore larger feature space and training data
- Cons
  - Modelling limit: feature independence assumption
  - Cannot capture feature interactions unless defining high order combination features
    - E.g., hour=10AM & city=London & browser=Chrome

# Non-linear Models

- ▶ Factorization Machines
- ▶ Gradient Boosting Decision Trees
- ▶ Combined Models
- ▶ Deep Neural Networks

# Factorization Machines

- Prediction based on feature embedding

$$y_{\text{FM}}(\mathbf{x}) := \text{sigmoid}\left(\underbrace{w_0 + \sum_{i=1}^N w_i x_i}_{\text{Logistic Regression}} + \underbrace{\sum_{i=1}^N \sum_{j=i+1}^N \langle \mathbf{v}_i, \mathbf{v}_j \rangle x_i x_j}_{\text{Feature Interactions}}\right)$$

For  $\mathbf{x} = [\text{Weekday=Friday}, \text{Gender=Male}, \text{City=Shanghai}]$

$$y_{\text{FM}}(\mathbf{x}) = \text{sigmoid}\left(w_0 + w_{\text{Friday}} + w_{\text{Male}} + w_{\text{Shanghai}} + \langle \mathbf{v}_{\text{Friday}}, \mathbf{v}_{\text{Male}} \rangle + \langle \mathbf{v}_{\text{Friday}}, \mathbf{v}_{\text{Shanghai}} \rangle + \langle \mathbf{v}_{\text{Male}}, \mathbf{v}_{\text{Shanghai}} \rangle\right)$$

[Rendle. Factorization machines. ICDM 2010.]

[Oentaryo et al. Predicting response in mobile advertising with hierarchical importance-aware factorization machine. WSDM 14]



# Field-aware Factorization Machines

- Feature embedding for another field

$$y_{\text{FFM}}(\mathbf{x}) = \text{sigmoid}\left(w_0 + \sum_{i=1}^N w_i + \underbrace{\sum_{i=1}^N \sum_{j=i+1}^N \langle \mathbf{v}_{i, \text{field}(j)}, \mathbf{v}_{j, \text{field}(i)} \rangle x_i x_j}_{\text{Field-aware field embedding}}\right)$$

For  $\mathbf{x}=[\text{Weekday=Friday}, \text{Gender=Male}, \text{City=Shanghai}]$

$$\begin{aligned} y_{\text{FFM}}(\mathbf{x}) = \text{sigmoid}\big( & w_0 + w_{\text{Friday}} + w_{\text{Male}} + w_{\text{Shanghai}} \\ & + \langle \mathbf{v}_{\text{Friday}, \text{Gender}}, \mathbf{v}_{\text{Male}, \text{Weekday}} \rangle + \langle \mathbf{v}_{\text{Friday}, \text{City}}, \mathbf{v}_{\text{Shanghai}, \text{Weekday}} \rangle \\ & + \langle \mathbf{v}_{\text{Male}, \text{City}}, \mathbf{v}_{\text{Shanghai}, \text{Gender}} \rangle \big) \end{aligned}$$

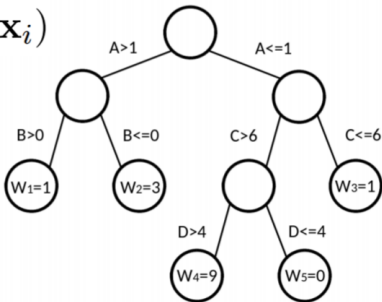
[Juan et al. Field-aware Factorization Machines for CTR Prediction. RecSys 2016.]

# Gradient Boosting

- Additive decision trees for prediction

$$\hat{y}_i = \phi(\mathbf{x}_i) = \sum_{k=1}^K f_k(\mathbf{x}_i), \quad f_k \in \mathcal{F}$$

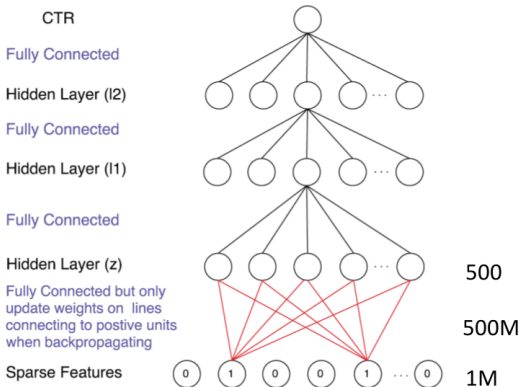
- Each decision tree  $f_k(\mathbf{x}_i)$



[Chen and He. Higgs Boson Discovery with Boosted Trees . HEPML 2014.]

# Neural Networks Models

- Difficulty:  
Impossible to  
directly deploy  
neural network  
models on such  
data



E.g., input features 1M, first layer 500, then 500M parameters for first layer

# Review Factorization Machines

- Prediction based on feature embedding

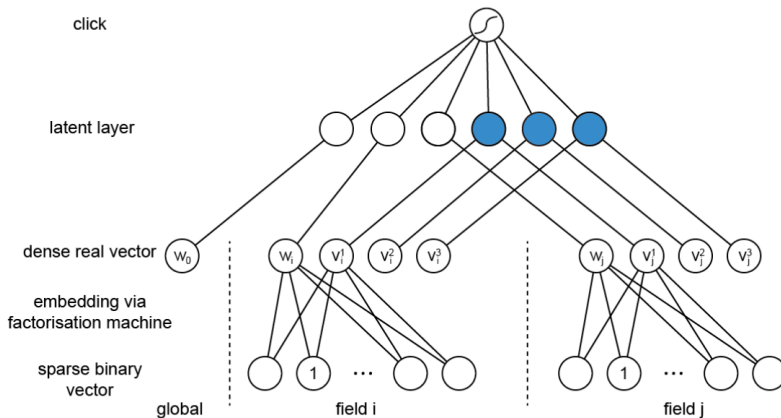
$$y_{\text{FM}}(\mathbf{x}) := \text{sigmoid} \left( \underbrace{w_0 + \sum_{i=1}^N w_i x_i}_{\text{Logistic Regression}} + \underbrace{\sum_{i=1}^N \sum_{j=i+1}^N \langle \mathbf{v}_i, \mathbf{v}_j \rangle x_i x_j}_{\text{Feature Interactions}} \right)$$

- Embed features into a k-dimensional latent space
- Explore the feature interaction patterns using vector inner-product

[Rendle. Factorization machines. ICDM 2010.]

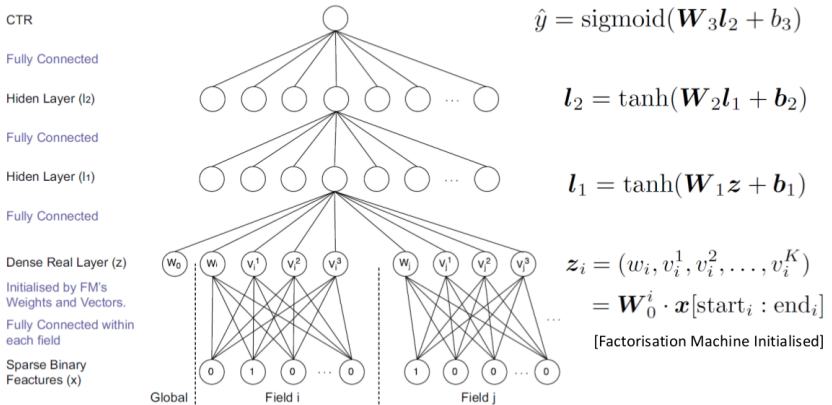
[Oentaryo et al. Predicting response in mobile advertising with hierarchical importance-aware factorization machine. WSDM 14]

# Factorization Machines is a Neural Networks



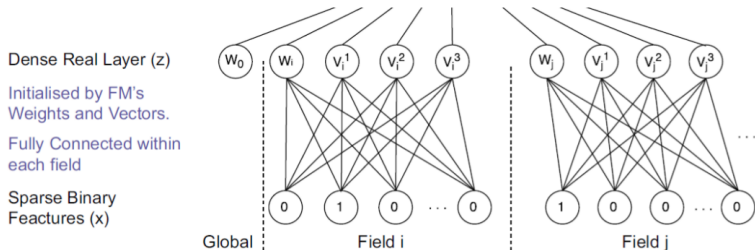
$$y_{\text{FM}}(\mathbf{x}) := \text{sigmoid}\left(w_0 + \sum_{i=1}^N w_i x_i + \sum_{i=1}^N \sum_{j=i+1}^N \langle \mathbf{v}_i, \mathbf{v}_j \rangle x_i x_j\right)$$

# Factorization-Machine supported Neural Networks(FNN)



[Zhang et al. Deep Learning over Multi-field Categorical Data – A Case Study on User Response Prediction. ECIR 16]

# Factorization-Machine supported Neural Networks(FNN)



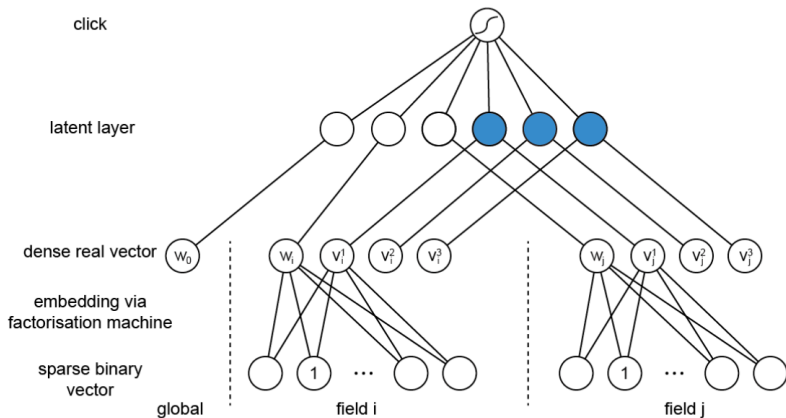
- Chain rule to update factorisation machine parameters

$$\frac{\partial L(y, \hat{y})}{\partial \mathbf{W}_0^i} = \frac{\partial L(y, \hat{y})}{\partial \mathbf{z}_i} \frac{\partial \mathbf{z}_i}{\partial \mathbf{W}_0^i} = \frac{\partial L(y, \hat{y})}{\partial \mathbf{z}_i} \mathbf{x}[\text{start}_i : \text{end}_i]$$

$$\mathbf{W}_0^i \leftarrow \mathbf{W}_0^i - \eta \cdot \frac{\partial L(y, \hat{y})}{\partial \mathbf{z}_i} \mathbf{x}[\text{start}_i : \text{end}_i].$$

[Zhang et al. Deep Learning over Multi-field Categorical Data – A Case Study on User Response Prediction. ECIR 16]

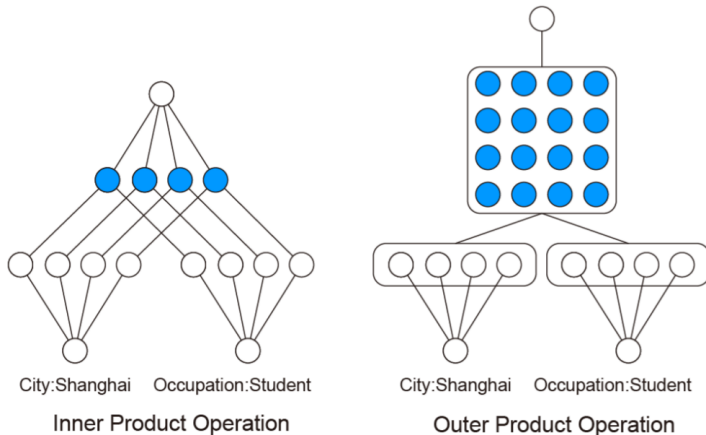
# Factorization-Machine different from Neural Networks



$$y_{\text{FM}}(\mathbf{x}) := \text{sigmoid}\left(w_0 + \sum_{i=1}^N w_i x_i + \sum_{i=1}^N \sum_{j=i+1}^N \langle \mathbf{v}_i, \mathbf{v}_j \rangle x_i x_j\right)$$

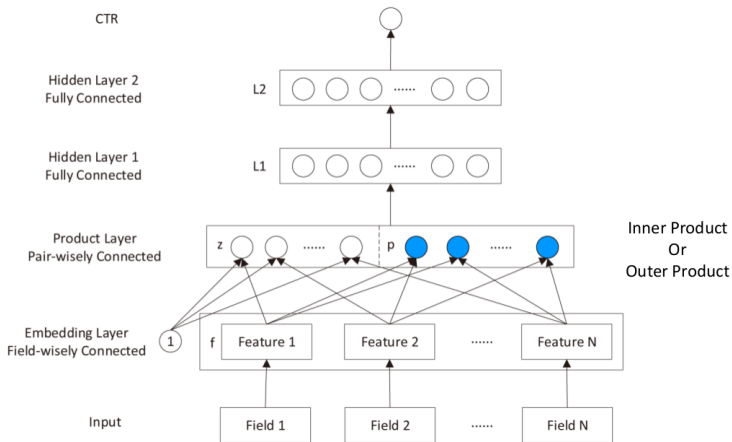


# Product Operations as Feature Interactions



[Yanru Qu et al. Product-based Neural Networks for User Response Prediction. ICDM 2016]

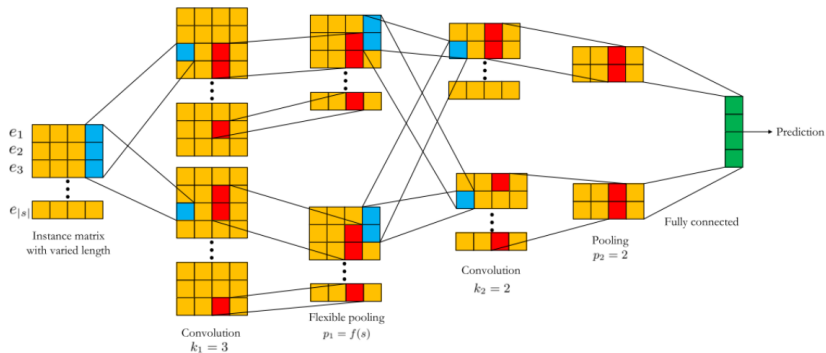
# Product-based Neural Networks (PNN)



[Yanru Qu et al. Product-based Neural Networks for User Response Prediction. ICDM 2016]

# Convolutional Click Prediction Model

- CNN to (partially) select good feature combinations



[Qiang Liu et al. A convolutional click prediction model. CIKM 2015]

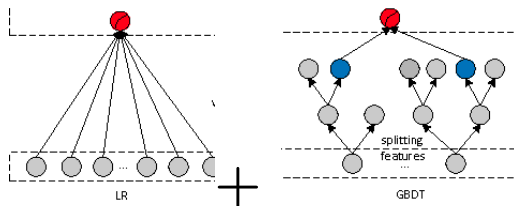
# Overall comparisons

Model	AUC		Log Loss	
	Criteo	iPinYou	Criteo	iPinYou
LR	71.48%	73.43%	0.1334	5.581e-3
FM	72.20%	75.52%	0.1324	5.504e-3
FNN	75.66%	76.19%	0.1283	5.443e-3
CCPM	76.71%	76.38%	0.1269	5.522e-3
PNN-I	<b>77.79%</b>	79.14%	<b>0.1252</b>	5.195e-3
PNN-II	77.54%	<b>81.74%</b>	0.1257	5.211e-3
PNN-III	77.00%	76.61%	0.1270	<b>4.975e-3</b>

Model	RMSE		RIG	
	Criteo	iPinYou	Criteo	iPinYou
LR	9.362e-4	5.350e-07	6.680e-2	7.353e-2
FM	9.284e-4	5.343e-07	7.436e-2	8.635e-2
FNN	9.030e-4	5.285e-07	1.024e-1	9.635e-2
CCPM	8.938e-4	5.343e-07	1.124e-1	8.335e-2
PNN-I	<b>8.803e-4</b>	4.851e-07	<b>1.243e-1</b>	1.376e-1
PNN-II	8.846e-4	5.293e-07	1.211e-1	1.349e-1
PNN-III	8.988e-4	<b>4.819e-07</b>	1.118e-1	<b>1.740e-1</b>

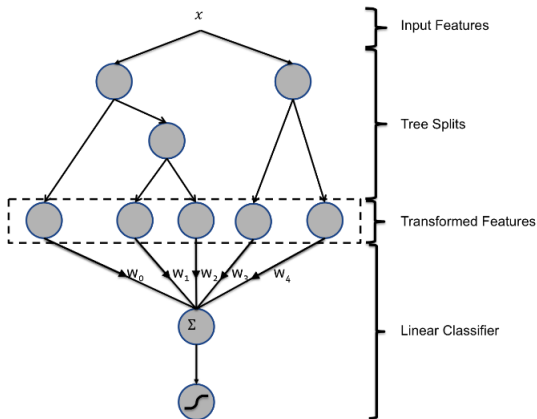
# Practical Lessons From Industry Companies

## Бустинг логрегрессии деревьями

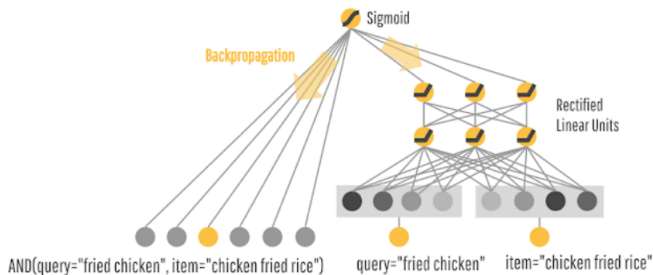


# Facebook (2014)

## Дообучаем деревья логрегрессий



## Совместно обучаем логрегрессию и нейронку





# Microsoft Bing (2017)

## Бустим нейронку деревьями

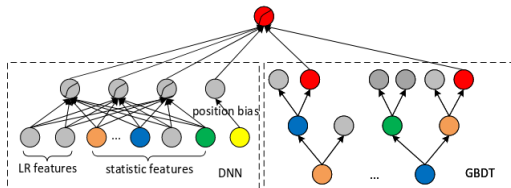


Figure 3: DNN+GBDT.

# Microsoft Bing (2017)

Models	Position=ML1		Position=ALL		Description
	AUC Gain	RIG Gain	AUC Gain	RIG Gain	
NN	0.00%	0.00%	0.00%	0.00%	NN with 1 hidden layer and 30 hidden units (baseline model)
LR	-1.97%	-16.14%	-1.46%	-10.01%	LR with normalized position bias
LR V2	-1.81%	-10.68%	-0.91%	-5.13%	LR with inversed position bias
GBDT2LR	0.06%	-0.17%	0.05%	0.44%	Cascade leaf index in GBDT as categorical feature to LR (used in Facebook [25])
LR+GBDT	0.12%	-1.87%	-0.33%	-1.93%	Boost LR with GBDT (used in Yandex [25])
LR2GBDT V2	0.13%	-0.14%	0.03%	0.67%	Cascade LR with inversed position bias to GBDT
GBDT	0.14%	0.36%	0.03%	0.91%	GBDT initialized with inversed position bias
LR2GBDT	0.14%	-0.27%	0.01%	0.50%	Cascade LR with normalized position bias to GBDT
GBDT2NN	0.16%	1.29%	0.04%	1.32%	Cascade GBDT to NN
LR+GBDT V2	0.24%	1.36%	0.07%	1.04%	Boost LR (inversed position bias) with GBDT
NN2GBDT	0.25%	0.15%	0.08%	0.72%	Cascade NN to GBDT
GBDT+DNN	0.25%	1.33%	0.15%	1.52%	Average NN and GBDT
NN+GBDT	0.40%	2.81%	0.15%	1.30%	Boost NN with GBDT

# Вопросы

