



ТЕХНОСФЕРА

Learning to match 2

Владимир Гулин

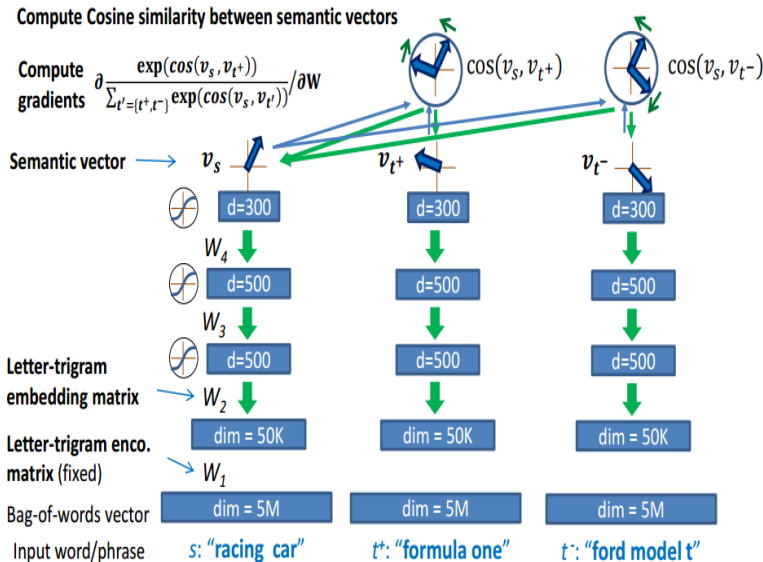
27 ноября 2019 г.

План лекции

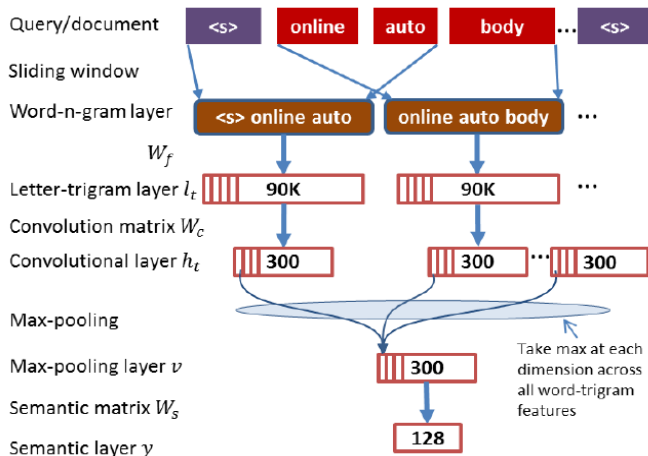
Напоминание

Обзор современных статей

DSSM



CLSM



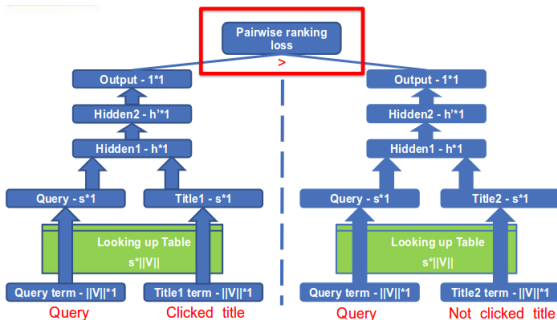
Вопрос:

- Как такие модели обучать?

Industry


Baidu DNN model

- ▶ Порядка 30% в улучшении качества за 2013 и 2014
- ▶ Используют 10 млрд. кликов для обучения (более 100 млн. параметров)



ARC-I

- ▶ Similar to DSSM, perform 1D convolution on text representations separately.
- ▶ Originally introduced for paraphrasing task.

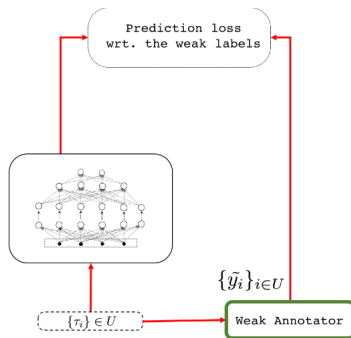


Training neural networks using pseudo relevance

Main idea


Размечаем большое количество
неразмеченных данных
используя слабую модель
(Pseudo-Labeling) и тренируем
модель на таком слабом
таргете

- ▶ Аппроксимируем BM25
- ▶ Побеждаем BM25
обучившись на BM25!



Training neural networks using pseudo relevance

- ▶ Используются 3 различные архитектуры Score, Rank, RankProb



- ▶ Моделям скормливаются три различных представления

- ▶ Dense: $\psi(q, d) = [N || avg(l_d)_D || l_d || \{df(t_i) || tf(t_i, d)\}_{1 \leq i \leq k}]$
- ▶ Sparse: $\psi(q, d) = [tfv_c || tfv_q || tfv_d]$
- ▶ Embedding-based: $\psi(q, d) = [\odot_{i=1}^{|q|} (\mathcal{E}(t_i^q), \mathcal{W}(t_i^q)) || \odot_{i=1}^{|d|} (\mathcal{E}(t_i^d), \mathcal{W}(t_i^d))]$

Training neural networks using pseudo relevance

Method	Robust04			ClueWeb		
	MAP	P@20	nDCG@20	MAP	P@20	nDCG@20
BM25	0.2503	0.3569	0.4102	0.1021	0.2418	0.2070
Score + Dense	0.1961 [▽]	0.2787 [▽]	0.3260 [▽]	0.0689 [▽]	0.1518 [▽]	0.1430 [▽]
Score + Sparse	0.2141 [▽]	0.3180 [▽]	0.3604 [▽]	0.0701 [▽]	0.1889 [▽]	0.1495 [▽]
Score + Embed	0.2423 [▽]	0.3501	0.3999	0.1002	0.2513	0.2130
Rank + Dense	0.1940 [▽]	0.2830 [▽]	0.3317 [▽]	0.0622 [▽]	0.1516 [▽]	0.1383 [▽]
Rank + Sparse	0.2213 [▽]	0.3216 [▽]	0.3628 [▽]	0.0776 [▽]	0.1989 [▽]	0.1816 [▽]
Rank + Embed	0.2811[▲]	0.3773[▲]	0.4302[▲]	0.1306[▲]	0.2839[▲]	0.2216[▲]
RankProb + Dense	0.2192 [▽]	0.2966 [▽]	0.3278 [▽]	0.0702 [▽]	0.1711 [▽]	0.1506 [▽]
RankProb + Sparse	0.2246 [▽]	0.3250 [▽]	0.3763 [▽]	0.0894 [▽]	0.2109 [▽]	0.1916
RankProb + Embed	0.2837[▲]	0.3802[▲]	0.4389[▲]	0.1387[▲]	0.2967[▲]	0.2330[▲]

Interaction matrix based approaches


Alternative to Siamese networks

Interaction matrix X , where $x_{i,j}$ is obtained by comparing the i^{th} word in source sentence with j^{th} word in target sentence

Comparisons can be both lexical or semantic


Вопрос:

- Как такое использовать с линейными моделями?




ARC-II

- Let two sentences meet **before** their own high-level representations mature
- Basic matching signals: phrase sum interaction matrix
- Interaction: CNN to capture the local interaction structure
- Aggregation Function: MLP



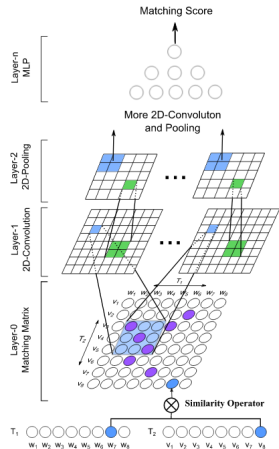
ARC-II

- Keeping word order information
 - Both the convolution and pooling are order preserving




- However, word level exact matching signals are lost
 - 2-D matching matrix is constructed based on the embedding of the words in two N-grams

MatchPyramid




MatchPyramid

Interaction matrix for two texts captures pairwise similarity between their words

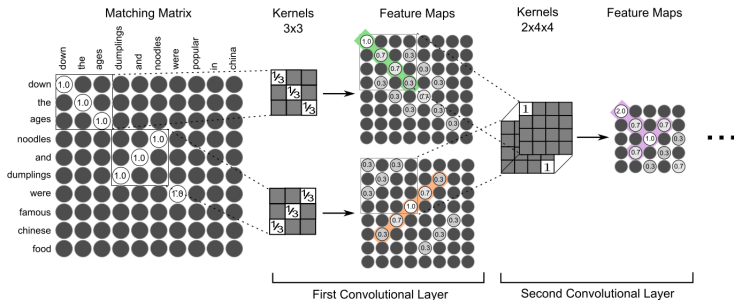


Exact match



Cosine similarity


Interaction Matrix: MatchPyramid



MatchPyramid

T₁: PCCW's chief operating officer, Mike Butcher, and Alex Arena, the chief financial officer, will report directly to Mr So.

T₂: Current Chief Operating Officer Mike Butcher and Group Chief Financial Officer Alex Arena will report to So.



MatchPyramid

Table 1: Results on MSRP.

Model	Acc.(%)	F ₁ (%)
ALLPOSITIVE	66.50	79.87
TF-IDF	70.31	77.62
DSSM	70.09	80.96
CDSSM	69.80	80.42
ARC-I	69.60	80.27
ARC-II	69.90	80.91
MP-IND	75.77	82.66
MP-Cos	75.13	82.45
MP-DOT	75.94	83.01

Table 2: Results on the task of paper citation matching.

Model	Acc.(%)	F ₁ (%)
ALLPOSITIVE	33.33	50.00
TF-IDF	82.63	70.21
DSSM	71.97	29.88
CDSSM	69.84	19.97
ARC-I	84.51	76.79
ARC-II	86.48	79.57
MP-IND	73.76	44.71
MP-Cos	86.65	79.70
MP-DOT	88.73	82.86

Lexical and Semantic matching

Query: united states president

The **President** of the **United States** of America (POTUS) is the elected head of state and head of government of the **United States**. The **president** leads the executive branch of the federal government and is the commander in chief of the United States Armed Forces. Barack Hussein Obama II (born August 4, 1961) is an American politician who is the 44th and current **President** of the United States. He is the first African American to hold the office and the first **president** born outside the continental **United States**.

The President of the **United States** of America (POTUS) is the elected head **of** state and head of government of **the United States**. The **president** leads the executive branch **of the federal government** **and is** the commander in chief of the United States Armed Forces. **Barack Hussein Obama II** (born August 4, 1961) is an American politician **who is** the 44th and current President of **the United States**. **He** is the first African American to hold **the** office and the first president born outside **the continental** United States.

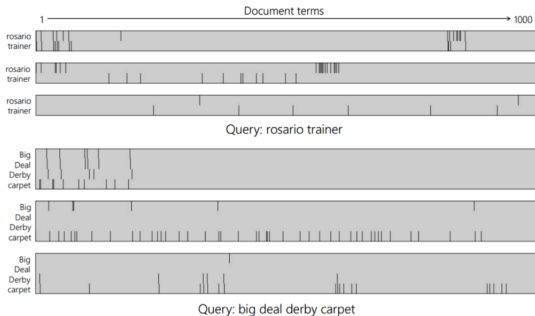
- ▶ Традиционные модели инфопоиска работают исключительно на лексическом матчинге
- ▶ Representation модели могут учитывать как все термины документа соотносятся с запросом
- ▶ И те и другие могут быть смоделированы с помощью нейросетей

Lexical Matching

Query: “rosario trainer”

The rare term “rosario” may have never been seen during training and unlikely to have meaningful representation

But the patterns of lexical matches of rare terms in document may be very informative for estimating relevance



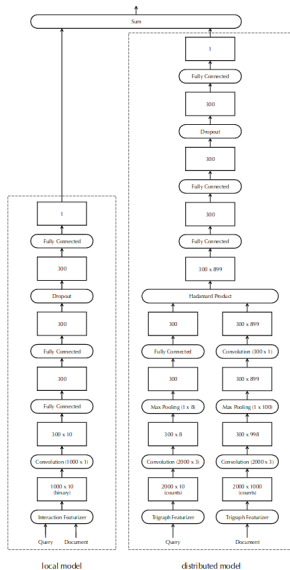
Duet Model

Совместно тренируем 2 подсетки, сфокусированных на лексическом и семантическом матчинге


Training sample:

$q, d^+, d_1, d_2, d_3, d_4$

$$p(d^+|q) = \frac{e^{nrdm(q,d^+)}}{\sum_{d \in D^-} e^{nrdm(q,d)}}$$



Distributed Model



Duet Model

(a) weighted

	NDCG@1	NDCG@10
Non-neural baselines		
LSA	22.4	44.2
BM25	24.2	45.5
DM	24.7	46.2
QL	24.6	46.3
Neural baselines		
DRMM	24.3	45.2
DSSM	25.8	48.2
CDSSM	27.3	48.2
DESM	25.4	48.3
Our models		
Local model	24.6	45.1
Distributed model	28.6	50.5
Duet model	32.2	53.0

(b) unweighted

	NDCG@1	NDCG@10
Non-neural baselines		
LSA	31.9	62.7
BM25	34.9	63.3
DM	35.0	63.4
QL	34.9	63.4
Neural baselines		
DRMM	35.6	65.1
DSSM	34.3	64.4
CDSSM	34.3	64.0
DESM	35.0	64.7
Our models		
Local model	35.0	64.4
Distributed model	35.2	64.9
Duet model	37.8	66.4

Duet Model

Term importance



Local model

Only query terms have an impact

Earlier occurrences have bigger impact

Query: united states president

The **President** of the **United States** of America (POTUS) is the elected head of state and head of government of the **United States**. The **president** leads the executive branch of the federal government and is the commander in chief of the United **States** Armed Forces. Barack Hussein Obama II (born August 4, 1961) is an American politician who is the 44th and current **President** of the United States. He is the first African American to hold the office and the first **president** born outside the continental **United States**.

Visualizing impact of dropping terms on model score

Duet Model

Term importance

Distributed model

Non-query terms (e.g., *Obama* and *federal*) has positive impact on score

Common words like 'the' and 'of' probably good indicators of well-formedness of content

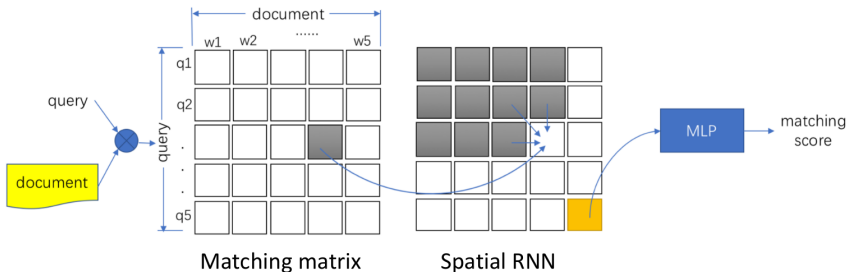
Query: united states president

The President of the United States of America (POTUS) is the elected head of state and head of government of the United States. The president leads the executive branch of the federal government and is the commander in chief of the United States Armed Forces. Barack Hussein Obama II (born August 4, 1961) is an American politician who is the 44th and current President of the United States. He is the first African American to hold the office and the first president born outside the continental United States.



Visualizing impact of dropping terms on model score

Match-SRNN

- Based on spatial recurrent neural network (SRNN)
- Basic matching signals: word-level matching matrix
- Matching function: Spatial RNN + MLP



Match-SRNN




- Calculated recursively (from top left to bottom right)
- All matching signals between the prefixes been utilized
 - **Current position:** sat <—> balls
 - **Substrings:**
 - the cat <—> the dog play
 - the cat <—> the dog play balls
 - the cat sat <—> the dog play

Match-SRNN


*Question: “How to get rid of **memory stick error** of my sony cyber shot?”*

*Answer: “You might want to try to format the **memory stick** but what is the **error** message you are receiving.”*



K-NRM

- ▶ Compute word-interaction matrix, apply k kernels to every query term row in interaction matrix.
- ▶ This results in k -dimensional vector.
- ▶ Aggregate the query term vectors into a fixed-dimensional query representation.




DeepRank

Query : Hubble Telescope Achievements

Document

Query-Centric Assumption



measurements the Hubble Space Telescope astronomers think they have ✓

using the Earth-orbiting Hubble Telescope measuring distances to 800 stars ✓

since 1929 when Edwin P. Hubble discovered that the galaxies are ✗



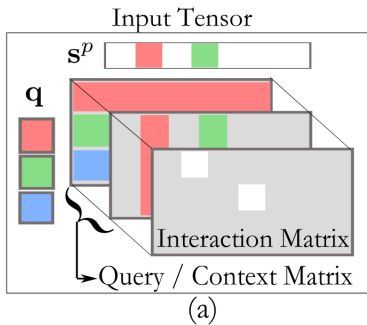
expansion rate also known as Hubble constant are fundamental to all ✗

expansion rate using Earth-orbiting Hubble telescope in measuring distances ✓

...


} Relevance ?

DeepRank



DeepRank

- ▶ Focus only on exact term occurrences in document.
- ▶ Compute interaction between query and window surrounding term occurrence.
- ▶ RNN or CNN then combines per-window features (query representation, context representations and interaction between query/document term) into matching score.



DeepRank

MQ2008									
Model	NDCG@1	NDCG@3	NDCG@5	NDCG@10	P@1	P@3	P@5	P@10	MAP
BM25-TITLE	0.344 ⁻	0.420 ⁻	0.461 ⁻	0.220 ⁻	0.408 ⁻	0.381 ⁻	0.337 ⁻	0.245 ⁻	0.465 ⁻
RANKSVM	0.375 ⁻	0.431 ⁻	0.479 ⁻	0.229	0.441 ⁻	0.390 ⁻	0.348 ⁻	0.249	0.478 ⁻
RANKBOOST	0.381	0.436 ⁻	0.477 ⁻	0.231	0.455	0.392 ⁻	0.347 ⁻	0.248	0.481 ⁻
ADARANK	0.360 ⁻	0.422 ⁻	0.462 ⁻	0.222	0.430 ⁻	0.384 ⁻	0.339 ⁻	0.247 ⁻	0.468 ⁻
LAMBDA MART	0.378	0.437 ⁻	0.477 ⁻	0.231	0.446 ⁻	0.398	0.348 ⁻	0.251	0.478 ⁻
DSSM	0.286 ⁻	0.336 ⁻	0.378 ⁻	0.178 ⁻	0.341 ⁻	0.307 ⁻	0.284 ⁻	0.221 ⁻	0.391 ⁻
CDSSM	0.283 ⁻	0.331 ⁻	0.376 ⁻	0.175 ⁻	0.335 ⁻	0.302 ⁻	0.279 ⁻	0.222 ⁻	0.395 ⁻
ARC-I	0.295 ⁻	0.363 ⁻	0.413 ⁻	0.187 ⁻	0.361 ⁻	0.336 ⁻	0.311 ⁻	0.229 ⁻	0.424 ⁻
SQA-NOFEAT	0.291 ⁻	0.350 ⁻	0.401 ⁻	0.184 ⁻	0.366 ⁻	0.332 ⁻	0.309 ⁻	0.231 ⁻	0.416 ⁻
DRMM	0.368 ⁻	0.427 ⁻	0.468 ⁻	0.220 ⁻	0.437 ⁻	0.392 ⁻	0.344 ⁻	0.245 ⁻	0.473 ⁻
ARC-II	0.299 ⁻	0.340 ⁻	0.394 ⁻	0.181 ⁻	0.366 ⁻	0.326 ⁻	0.305 ⁻	0.229 ⁻	0.413 ⁻
MATCHPYRAMID	0.351 ⁻	0.401 ⁻	0.442 ⁻	0.211 ⁻	0.408 ⁻	0.365 ⁻	0.329 ⁻	0.239 ⁻	0.449 ⁻
MATCH-SRNN	0.369 ⁻	0.426 ⁻	0.465 ⁻	0.223 ⁻	0.432 ⁻	0.383 ⁻	0.335 ⁻	0.239 ⁻	0.466 ⁻
DEEPRANK-2DGRU	0.391	0.436	0.480	0.236	0.462	0.395	0.354	0.252	0.489
DEEPRANK-CNN	0.406	0.460	0.496	0.240	0.482	0.412	0.359	0.252	0.498
SQA	0.402	0.454	0.493	0.236	0.485	0.411	0.362	0.254	0.496
DEEPRANK-CNN-FEAT	0.418	0.475	0.507	0.248	0.497	0.422	0.366	0.255	0.508


Связь с классикой

$$\text{BM25}(\mathbf{q}, \mathbf{d}) = \sum_{w \in \mathbf{q}} \text{IDF}(w) \cdot \frac{f(w, \mathbf{d}) \cdot (k_1 + 1)}{f(w, \mathbf{d}) + k_1 \cdot (1 - b + \frac{b|\mathbf{d}|}{\text{avgdl}})}$$


$$\begin{aligned} \text{DeepRank}(\mathbf{q}, \mathbf{d}) &= \sum_{w \in \mathbf{q}} \text{IDF}(w) \cdot \text{RNN}_{p \in \mathbb{P}(w)} \left[1, 1/|\mathbf{d}| \right] \\ &= \sum_{w \in \mathbf{q}} \text{IDF}(w) \cdot \mathcal{G}(f(w, \mathbf{d}), |\mathbf{d}|), \end{aligned}$$

Summary

- Representation learning:
representing queries and
document in
semantic space



- Matching function learning:
discovering and aggregating
the query-document
matching patterns



Вопросы

